

Ain Foundry Virtual EIR

Product and Operating Blueprint for Persistent, Human-Led Venture Coaching



1:1 persistent vEIR-to-venture assignment	24/7 practice, diagnosis, and preparation	Human-led escalation and final authority
---	---	--

Prepared by	Satvik Sathish Adyanthaya
Project sponsor	Professor Roberto Colangelo, Executive Director, Ain Center for Entrepreneurship and Innovation
Audience	Ain Center leadership, Entrepreneurs-in-Residence, program operators, product owners, and implementation partners
Date	September 2026

Working name: Ain Foundry Virtual Entrepreneur-in-Residence (vEIR). This document proposes a human-led, AI-amplified capability; it does not authorize deployment or the use of any person's likeness.

OPERATING INTENT

1. Product thesis and north star

Product thesis

Ain Foundry can multiply the reach of its EIR network by giving every venture a persistent, synthetic coaching interface that retains approved context, challenges assumptions, teaches venture-building practice, and prepares better human conversations.

The platform is not a digital spokesperson, a generic entrepreneurship chatbot, or an autonomous advisor. It is a governed interface around four durable assets: **the Ain venture-building method, each EIR's approved doctrine, each venture's private operating memory, and an explicit human escalation network.** The avatar makes the interaction natural; the defensible product value sits in continuity, evidence, program integration, and trust. [S1]

~120	7	~17	4
venture-scale planning baseline	EIR profiles currently listed [S3]	ventures per listed profile if evenly distributed	north-star dimensions: access, quality, continuity, safety

North-star metric

Qualified venture progress per scarce human EIR hour. Qualified progress means an evidence-backed decision, validated or retired assumption, completed milestone, high-quality external interaction, or correctly routed escalation - not merely a conversation count.

Positioning

What the vEIR does Diagnoses, challenges, teaches, rehearses, retrieves approved resources, maintains continuity, and prepares structured escalations.	What remains human Consequential judgment, relationship-building, reputation-bearing advice, exceptions, sensitive conflicts, external commitments, and program governance.
---	--

Product principles

- **Continuity before charisma:** accurate venture memory matters more than visual spectacle.
- **Evidence before eloquence:** a persuasive unsupported answer is a failure.
- **Challenge before comfort:** the system seeks disconfirming evidence and names uncertainty.
- **Program-native:** advice should route into Ain mentors, UR Ventures, competitions, labs, funding pathways, and external incubators where appropriate.
- **Consent and reversibility:** an EIR can inspect, correct, suspend, and retire the persona.

ROUTE TO THE RIGHT HUMAN LAYER

1A. EIR network and Foundry service model




Figure 1. Seven profiles currently appear on the official Ain EIR directory; expertise labels summarize that public directory and must be validated before persona or assignment decisions. [S3]

Human layer	Officially described contribution	vEIR operating role
Entrepreneurs-in-Residence	Seasoned professionals providing strategic guidance, mentorship, idea refinement, and challenge navigation. [S3]	Define doctrine; approve persona behavior; receive judgment-heavy escalations; calibrate quality.
Foundry Consultants	MBA, MS, and PhD students supporting undergraduate teams weekly on customer discovery, business models, and planning. [S17]	Support recurring work plans, evidence collection, and structured follow-through around vEIR sessions.
Foundry Mentors	Alumni and practitioners working with one or two teams, workshops, feedback, introductions, and ecosystem access. [S18]	Receive relationship-bearing referrals, domain introductions, pitch feedback, and scale-stage coaching.

Routing principle

The vEIR prepares and routes; it does not flatten EIRs, consultants, and mentors into interchangeable avatars. Human role, assignment, consent, availability, and expertise determine the destination.

EXPERIENCE ARCHITECTURE

2. Users, roles, and jobs to be done

Role	Primary jobs	Product promise	Authority boundary
Founder / venture team	Clarify a decision, prepare for an interaction, test assumptions, find a resource, maintain momentum.	Context persists; guidance is available when work happens; outputs are actionable and source-linked.	Founder owns venture choices and can correct eligible memory.
Human EIR	Scale coaching, preserve doctrine, inspect quality, handle high-value escalations.	Less context reconstruction and more time on judgment-intensive work.	Approves persona behavior and remains final advisory authority.
Ain program lead	Allocate resources, see portfolio needs, govern quality, coordinate pathways.	Program signals emerge without exposing raw venture secrets.	Controls enrollment, policies, routing, and suspension.
Domain specialist	Answer a narrow escalated question with adequate context.	Receives a concise, authorized packet instead of a vague referral.	Sees only the information necessary for the referral.
Security / privacy / accessibility reviewer	Verify controls and investigate incidents.	Auditable consent, access, data movement, and model behavior.	Can block or constrain deployment.

Founder jobs by moment

Moment	Founder asks	vEIR behavior	Expected output
Idea formation	Is this a problem worth pursuing?	Separates problem, customer, evidence, and assumption; demands a falsifiable next test.	Problem statement and discovery plan.
Customer discovery	What did these interviews actually prove?	Challenges confirmation bias and compares evidence against the hypothesis.	Evidence ledger and revised hypothesis.
Business model	Can this become viable?	Tests willingness to pay, unit logic, channels, constraints, and alternatives.	Business-model risks and experiments.
Commercialization	What path fits this technology?	Routes general process questions and flags IP, licensing, FDA, reimbursement, or lab constraints.	Pathway map plus specialist escalation.
Pitch / sales / board	Can I practice this conversation?	Role-plays, interrupts weak logic, scores evidence, and provides a focused replay.	Rubric, corrections, and next rehearsal.
Transition	What should happen after Ain?	Matches readiness and needs to competitions, funders, programs, or incubators.	Readiness evidence and referral packet.

END-TO-END EXPERIENCE

3. Founder journey and service promise




Figure 2. The vEIR accompanies the venture from onboarding through external transition; human intervention is available at every phase.

Session contract

Every substantive vEIR session should end with five visible objects:

Object	Definition	Founder control
Diagnosis	Current stage, immediate constraint, and the assumption most likely to invalidate the plan.	Accept, revise, or reject.
Evidence	Sources and venture facts used, with conflicts or missing information made explicit.	Open, correct, or flag.
Challenge	The strongest skeptical question or counter-hypothesis the founder must address.	Respond now or save.
Next action	One or more specific experiments, decisions, or preparation steps with an owner and target date.	Edit, assign, or dismiss.
Escalation	Whether a human is required, why, what information may be shared, and the recommended destination.	Approve the handoff packet.

Interaction modes

- **Diagnose:** identify stage, constraint, risk, and missing evidence.
- **Challenge:** red-team assumptions, surface alternatives, and resist agreeable but weak conclusions.
- **Coach:** use questions and reflection to build founder judgment rather than completing all thinking for the founder.
- **Teach:** explain frameworks with examples tied to the venture's current decision.
- **Rehearse:** simulate investor, customer, board, partner, or stakeholder interactions and score the attempt.
- **Route:** match needs to Ain, University, regional, regulatory, funding, or specialist resources.

WHAT MUST EXIST

4. Capability blueprint




Figure 3. Product capabilities. The avatar is one interface to the operating system, not the operating system itself.

Capability releases

Capability	Capstone V1	Controlled pilot	Foundry scale
Identity and access	Invite-only pilot accounts; venture role and EIR assignment.	UR SSO if approved; delegated admin; join / leave lifecycle.	Program-wide provisioning, access review, and policy enforcement.
Conversation	Voice or text plus a real-time / near-real-time avatar session.	Interruptions, push-to-talk, captions, accessibility controls.	Concurrency management, service recovery, and multi-device experience.
Venture memory	Facts, decisions, artifacts, milestones, corrections.	Memory suggestions, approval workflow, source aging, conflict resolution.	Portfolio-safe aggregates and lifecycle retention.
EIR persona	One doctrine and style profile; representative scenario approval.	Two personas; calibration studio and versioning.	Active-persona registry; the seven currently listed profiles remain subject to consent, availability, recertification, suspension, and retirement.
Evidence and tools	Approved corpus retrieval and citations.	Curated web / program resources and pathway tools.	Connector governance, freshness SLAs, and source ownership.
Escalation	Structured packet and manual routing.	Workflow queues, consented sharing, status tracking.	Routing rules, service levels, portfolio analytics, and exception management.
Quality and safety	Scenario tests, red-team suite, audit events.	Continuous sampling, incident workflow, founder feedback.	Dashboards, model / vendor change control, and formal review cadence.

MINIMAL INTERFACE; VISIBLE EVIDENCE

4A. Founder workspace shell

Founder session shell — muted black, gray, and evidence-first




Figure 4. The session shell keeps the avatar, transcript, venture context, evidence, memory, and escalation controls in one restrained workspace.

Persistent shell	Nested workspace pages
Left navigation	Overview; Sessions; Memory; Evidence; Decisions; Experiments; Milestones; Documents; Team; Escalations
Top context bar	Current venture; stage; assigned EIR; synthetic persona release; session and consent status
Central work canvas	Avatar / text conversation; transcript; rehearsal target; current diagnostic or coaching mode
Right inspector	Claim sources; known facts; assumptions; contradictions; proposed memory; human handoff
Bottom control dock	Mic; camera; upload; stop; share; approve memory; escalate; keyboard and text alternatives

The application uses a muted black and gray system, large readable controls, clear hierarchy, restrained motion, and non-color status cues. The avatar may disappear in text-only or degraded mode without removing any core founder job.

MEMORY AS PRODUCT

5. Persistent venture workspace




Figure 5. Retrieval combines shared venture-building knowledge with EIR doctrine and venture-private memory; the current session is the most temporary layer.

Memory object model

Object	Examples	Update rule	Expiry / review
Venture fact	Customer segment, team role, product constraint.	Source-linked; founder or authorized user confirms material facts.	Review when source changes or becomes stale.
Hypothesis	Customer pain, channel, price, technical feasibility.	Versioned with supporting and contradicting evidence.	Closed as validated, rejected, or unresolved.
Decision	Choose segment, pause feature, enter competition.	Record decision owner, rationale, date, and reversal condition.	Retained for active program unless policy says otherwise.
Experiment	Interview set, landing page, prototype, pricing test.	Track prediction, method, result, interpretation, and next step.	Archived when complete; remains sourceable.
Milestone	MVP, pilot, IP review, grant, pitch event.	Owner, target, status, dependency, evidence.	Escalate when blocked or overdue by rule.
Session insight	Risk, open question, corrected assumption.	Proposed by AI; visible to user before promotion to durable memory.	Discard unless approved or used in a decision.

Memory quality rules

- No durable memory from an unverified model inference alone.
- Every memory entry carries source, owner, timestamp, confidence, visibility, and correction history.
- Contradictory evidence coexists until a human resolves it; the system must not silently overwrite inconvenient facts.
- The founder can see what the vEIR remembers and why it was retrieved.

GOVERNED FIDELITY

6. EIR persona and synthetic identity lifecycle




Figure 6. No EIR persona reaches founders without explicit consent, approved doctrine, representative testing, and a reversible lifecycle.

Four separable fidelity layers

Layer	What it captures	How to build	Failure if conflated
Expertise	Frameworks, cases, heuristics, domain knowledge, referrals.	Curated sources, structured interviews, annotated cases, EIR review.	A realistic face gives generic or incorrect advice.
Judgment	How the EIR weighs evidence, uncertainty, risk, and stage.	Decision scenarios, counterexamples, rubrics, escalation doctrine.	The system mimics phrases but cannot reason in the EIR's way.
Communication style	Questioning pattern, directness, pacing, vocabulary, teaching style.	Consented transcripts, style card, negative examples, blind scoring.	Stereotype or caricature replaces authentic behavior.
Embodiment	Voice, accent, facial appearance, lip-sync, listening motions.	Consented voice and video capture through an approved provider.	Deepfake risk, uncanny output, or reputation harm.

Persona card - required fields

- Approved name, title, disclosure language, audiences, use cases, and expiration date.
- Domain strengths, explicit limits, referral destinations, and topics requiring immediate escalation.
- Reasoning doctrine: questions typically asked, evidence standard, risk tolerance, teaching principles, and counterexamples.
- Communication profile: directness, warmth, pace, preferred vocabulary, feedback structure, and behaviors to avoid.
- Representative gold-standard sessions, unacceptable sessions, quality threshold, approver, current version, and suspension status.

Fidelity rule

The objective is not perfect imitation. The objective is faithful, consented, bounded behavior that the EIR recognizes as professionally aligned and can correct.

WHO DOES WHAT

7. Human-AI operating model




Figure 7. Founders interact continuously; human EIRs govern doctrine, receive high-value escalations, and calibrate the system.

Escalation service levels - proposed

Priority	Examples	System behavior	Human target
P0 - Security / identity	Cross-venture exposure, compromised account, likeness misuse.	Block path, preserve minimal audit evidence, suspend affected access.	Immediate administrator / security response.
P1 - Consequential	IP rights, clinical / FDA, legal commitment, securities, safety, crisis.	Explain boundary, collect minimum context, request founder consent to route.	One business day or program-defined faster path.
P2 - Expert judgment	Strategic pivot, pricing, fundraising readiness, partnership negotiation.	Provide options and evidence; prepare an EIR packet.	Next available EIR review window.
P3 - Routine coaching	Discovery design, narrative, milestone planning, framework explanation.	Handle directly; include evidence, uncertainty, and next action.	Human review by sampling or founder request.

Escalation packet

- Founder-stated question and desired decision, not a model-written reinterpretation alone.
- Relevant venture facts and artifacts with permission to share.
- Evidence already considered, conflicting signals, and explicit unknowns.
- Why the vEIR escalated, urgency, and requested human role.
- A compact session history and proposed next step, clearly labeled as AI-generated.

RUN THE SERVICE

8. Program operations and governance

Cadence	Activity	Owner	Artifact
Every session	Disclosure, authorization, grounding, safety checks, recap, memory proposal.	Platform with founder control.	Session record and approved memory updates.
Weekly	Review P0 / P1 events, low-rated sessions, unresolved escalations, and source failures.	Program operator plus technical owner.	Quality and incident queue.
Biweekly in pilot	EIR calibration against representative sessions and emerging founder needs.	Participating EIR.	Persona / doctrine change set.
Monthly	Access review, corpus freshness, usage and cost, vendor availability, accessibility feedback.	Product owner and reviewers.	Operating scorecard.
Quarterly	Persona recertification, risk assessment, policy review, retention and deletion audit.	Ain leadership and EIR approver.	Release authorization.
On material change	Re-test model, prompt, retrieval, voice, avatar, policy, or vendor behavior.	Technical owner with EIR adjudication.	Change record and regression evidence.

Accountability map

Control	Accountable	Responsible	Consulted
Program purpose and acceptable use	Ain leadership	Product owner	EIRs, University reviewers, founders
Persona accuracy and release	Named EIR	Persona curator	Ain lead, evaluator
Data access and retention	Authorized University owner	System administrator	Privacy, security, records stakeholders
Advice quality and escalation	Ain product / program owner	EIR reviewers and operators	Domain specialists
Technical reliability	Technical owner	Engineering / vendor team	Program operator
Incident response	Designated University / Ain authority	Technical and program responders	Affected EIRs and founders

Founder transparency

The interface should show the persona version, synthetic disclosure, source panel, memory used, uncertainty, escalation status, data-use summary, and a clear way to report a problem. NIST's GenAI profile supports lifecycle risk governance; University guidance requires data classification and approved handling of non-public information. [S11][S12][S13][S14]

SHARED INTELLIGENCE

9. Portfolio learning without portfolio leakage




Figure 8. Shared learning is a governed publishing process, not direct retrieval across venture workspaces.

Pattern publication workflow

1. Detect a candidate lesson from multiple events or a deliberately nominated case.
2. Remove venture, founder, customer, lab, sponsor, financial, and technical identifiers.
3. Convert the lesson into a general pattern: context, signal, intervention, outcome, limits, and source class.
4. Review for re-identification, IP, bias, and whether the pattern is actually transferable.
5. Approve and version it in the shared Ain knowledge plane with owner and expiry date.
6. Monitor later use and withdraw the pattern if it produces misleading or harmful advice.


Portfolio signals allowed for program operations

Safe aggregate	Use	Guardrail
Count of ventures blocked by stage or topic	Plan workshops, mentors, and office hours.	Minimum group-size threshold; no raw text.
Escalation demand by domain	Recruit or schedule specialists.	Only category, urgency, and status unless authorized.
Resource search failure rate	Improve the Ain resource graph.	Do not expose the underlying venture query.
Milestone movement	Understand program progress and bottlenecks.	Report aggregated stage change, not confidential details.
Quality and safety signals	Calibrate personas and controls.	Restricted reviewer access and pseudonymized records.

Default rule. Publish only when a reasonable reader could not identify the originating venture, founder, lab, technology, customer, or negotiation.

SCENARIO MODEL

10. Capacity and economics



Illustrative: 120 ventures, 20-minute sessions, and \$0.10/min Avatar Only rate. Excludes LLM, ASR, TTS, storage, support, and implementation.

Figure 9. Illustrative monthly usage at program scale. The body-layer estimate uses a public Avatar Only credit rate as a modeling input, not a vendor quote. [S7][S8]

Cost architecture

Cost block	Driver	Management lever
Avatar streaming	Session minutes, resolution, mode, and provider tier.	Push-to-talk, session caps, audio-only fallback, approved overage policy.
Speech	ASR and TTS seconds, voice model, languages, concurrency.	Streaming model choice, caching of fixed messages, silence detection.
Reasoning and retrieval	Tokens, context size, documents retrieved, tools, retries.	Model routing, memory summarization, retrieval limits, response budgets.
Data and security	Storage, vector index, logs, keys, backups, retention.	Data minimization, tiered retention, low-volume archival.
People and operations	EIR calibration, content curation, incident response, user support.	Sampling plans, clear doctrine, reusable evaluation suites.
Implementation	Integration, accessibility, procurement, security review, maintenance.	Provider abstraction, narrow V1, staged rollout, reproducible deployment.

Capacity interpretation

- One 20-minute session per venture per month equals 40 live-avatar hours; four sessions equal 160 hours.
- The economic claim is not that AI minutes are free. It is that routine, repeatable support can be supplied at a predictable marginal cost while human EIR time moves toward higher-consequence work.
- The pilot must measure substitution and preparation effects. Usage alone does not prove human capacity was created.
- Before contracting, validate concurrency, session caps, data retention, training use, likeness rights, accessibility, incident support, and exit / deletion obligations.

STAGE GATES

11. Rollout and change management




Figure 10. Proposed rollout moves from one governed persona and roughly 10 ventures to full Foundry scale only after explicit gates.


Stage	Purpose	Entry gate	Exit gate
Capstone prototype	Prove the core experience, memory, isolation, quality, and handoff.	Sponsor charter; one EIR path; safe data; prototype resources.	No critical leakage; acceptable EIR quality; documented limitations and handoff.
Controlled pilot	Test real operations with 2 EIR personas and about 30 ventures.	Institutional approvals; owner and operator; support and incident process.	Stable service, founder value, human leverage, manageable cost, and governance evidence.
Foundry scale	Serve approximately 120 ventures with the active EIR network.	Budget, service owner, production security, accessibility, vendor terms, and training.	Quarterly recertification, portfolio-safe analytics, and expansion decision.

Adoption design

- Introduce the vEIR through live sessions led by the real EIR so founders understand purpose and boundaries.
- Start with concrete jobs - discovery debrief, pitch rehearsal, milestone unblock - rather than an empty chat box.
- Make escalation a positive feature, not a failure state; show founders how the AI improves the human meeting.
- Give EIRs a calibration studio and time budget; adoption will fail if persona maintenance is invisible labor.
- Publish quality, privacy, and synthetic-identity disclosures in plain language and repeat them inside the experience.

MANAGE OUTCOMES

12. Measurement system and decision dashboard



Leading indicators measure safe use; lagging indicators measure venture progress and EIR leverage.

Figure 11. KPI tree begins with qualified venture progress per human EIR hour, then balances access, quality, continuity, and safety.

Lens	Leading indicators	Outcome indicators	Anti-metric
Access	Activation, off-hours use, time-to-first-guidance.	Coverage across founder types, stages, and disciplines.	Raw message volume.
Quality	Evidence open rate, EIR sample score, correction frequency.	Better founder preparation and decision clarity.	Visual realism alone.
Continuity	Memory approval, correction, and retrieval accuracy.	Less repeated context and faster follow-through.	Amount of data retained.
Safety	Escalation recall, denied access tests, incident response time.	No material leakage or deceptive identity events.	Low escalation count by itself.
Human leverage	Completeness of handoff packets and review time.	More EIR time on consequential judgment.	Replacing EIR interactions.
Venture progress	Experiment completion and milestone movement.	Validated learning, stronger external interactions, appropriate next-program routing.	Fundraising as the only success measure.

Executive dashboard questions

1. Is founder access expanding across time, stage, discipline, and background without creating unequal quality?
2. Does the vEIR remember the right things, allow correction, and reduce repeated context?
3. Are founders arriving at human EIR sessions better prepared, with clearer evidence and questions?
4. Do privacy, consent, and escalation gates remain intact after every material model or vendor change?
5. Is the cost per qualified progress event defensible compared with additional manual support?

WHY AIN CAN BE DISTINCT

13. Strategic differentiation and external signal

HBS Foundry shows that founders will engage with AI-supported venture building and instructor avatars inside a structured program. Its public materials emphasize guided bootcamps, powered by an AI platform, while reporting identifies HeyGen-created faculty avatars used in practice interactions. [S9][S10] This is market validation, not a product specification for Ain.

Dimension	Generic entrepreneurship chatbot	Course / simulation avatar	Proposed Ain vEIR
Primary unit	One conversation	One exercise or program module	One evolving venture workspace
Memory	Limited or user-managed	Program-session context	Persistent, source-linked venture history
Persona	Generic assistant	Instructor representation	Consented EIR doctrine, style, voice, and likeness
Institutional context	Public information	Course content and platform pathways	Ain resources, UR commercialization paths, EIR and specialist network
Challenge behavior	Often agreeable	Scenario-dependent feedback	Explicit skepticism, evidence tests, coaching, and red-team mode
Human handoff	Usually external to the product	Live course touchpoints	Structured EIR / specialist escalation with venture context
Confidentiality model	Account-level vendor controls	Program platform controls	Venture isolation plus governed de-identified learning

Ain-specific moat

- A curated graph of University and regional venture resources, including commercialization and regulated pathways.
- Longitudinal venture memory and milestone history that survives changing schedules and program touchpoints.
- EIR-calibrated judgment and teaching patterns, with human review and revocation.
- A controlled learning loop from de-identified program patterns without pooling raw venture secrets.
- Direct integration between AI preparation and the real people who make the Ain ecosystem valuable.

LEADERSHIP HANDOFF

14. Decisions and risks

Decisions required before pilot build

Decision	Recommended position	Owner
V1 identity	One named, consenting EIR; neutral surrogate body until likeness approval is complete.	EIR and Ain leadership
Product architecture	Own the vEIR brain, memory, evidence, and policy; place the avatar vendor behind an adapter.	Technical owner
Data boundary	Public, approved program, EIR, and synthetic / expressly authorized venture data only in capstone V1.	Authorized University data owner
Pilot cohort	8-12 ventures in capstone; about 30 only after the first gate.	Program owner
Evaluation	Expert quality plus privacy, grounding, continuity, accessibility, founder value, and human leverage.	Capstone / evaluation lead
Scale	No full Foundry launch until security, privacy, procurement, accessibility, operations, and vendor terms are approved.	Ain and University authorities

Top operating risks

- The avatar's realism can cause over-trust or imply personal endorsement by the EIR.
- Venture isolation can fail through retrieval filters, logs, analytics, exports, support access, or vendor retention - not only the main database.
- Persona drift can turn rigorous skepticism into generic politeness, harsh imitation, or confidently wrong advice.
- External knowledge can become stale, while regulated and commercialization pathways vary by venture and time.
- EIR calibration and program operations can become an unfunded burden if ownership and cadence are not explicit.

EVIDENCE BASIS

15. Source register

The blueprint separates sponsor and program evidence, official EIR and Foundry role descriptions, University and NIST governance guidance, vendor capability documentation, and an external market signal. Vendor capabilities, access, retention, prices, and contract terms require written validation before implementation.

S1 Ain Center for Entrepreneurship and Innovation. Analytics Projects course sponsor briefing with Professor Roberto Colangelo. Video recorded August 31, 2026. User-provided source.

S2 University of Rochester News Center. Student entrepreneurs learn business by starting one. July 29, 2026.
<https://www.rochester.edu/newscenter/student-entrepreneurs-become-business-founders-712582/>

S3 Ain Center for Entrepreneurship and Innovation. Entrepreneurs in Residence. Accessed September 2026. <https://www.rochester.edu/aincenter/entrepreneurs-in-residence/>

S4 Ain Center for Entrepreneurship and Innovation. Home and Ain Foundry Program. Accessed September 2026. <https://www.rochester.edu/aincenter/>

S5 Microsoft Learn. Limited Access - custom neural voice and custom text to speech avatar. Updated August 26, 2026.
<https://learn.microsoft.com/en-us/azure/foundry/responsible-ai/speech-service/text-to-speech/limited-access>

S6 Microsoft Learn. Custom text to speech avatar overview and real-time synthesis. Accessed September 2026.
<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/text-to-speech-avatar/what-is-custom-text-to-speech-avatar>

S7 LiveAvatar Documentation. LiveAvatar Overview. Accessed September 2026. <https://docs.liveavatar.com/>

S8 HeyGen Help Center. LiveAvatar FAQ. Updated 2026. <https://help.heygen.com/en/articles/12758866-liveavatar-faq>

S9 Harvard Business School Foundry. Startup Bootcamp. Accessed September 2026. <https://www.hbs.edu/foundry/bootcamps/startup-bootcamp>

S10 TechCrunch. Harvard's \$699 startup bootcamp offers AI avatars of its instructors. August 22, 2026.
<https://techcrunch.com/2026/08/22/harvards-699-startup-bootcamp-offers-ai-avatars-of-its-instructors/>

S11 University of Rochester. Artificial Intelligence Tools. Accessed September 2026. <https://tech.rochester.edu/artificial-intelligence-ai-tools/>

S12 University of Rochester. Responsible Use of Generative Artificial Intelligence in Research. Accessed September 2026.
<https://www.rochester.edu/university-research/about/committees/ai-research-committee/responsible-use-gen-ai/>

S13 University of Rochester. Data Security Classification Policy. Accessed September 2026. <https://www.rochester.edu/policies/policy/data-security-classifications/>

S14 National Institute of Standards and Technology. AI RMF Generative AI Profile, NIST AI 600-1. Updated April 8, 2026. <https://doi.org/10.6028/NIST.AI.600-1>

S15 Project research note supplied by Satvik Sathish Adyanthaya: Azure virtual human, real-time EIR avatars, and HBS Foundry comparison. September 2026. User-provided source; claims independently checked where used.

S16 Ain Center for Entrepreneurship and Innovation. Ain Foundry Program. Accessed September 2026. <https://www.rochester.edu/aincenter/ain-foundry-program/>

S17 Ain Center for Entrepreneurship and Innovation. Foundry Consultants. Accessed September 2026. <https://www.rochester.edu/aincenter/foundry-consultants/>

S18 Ain Center for Entrepreneurship and Innovation. Foundry Mentors. Accessed September 2026. <https://www.rochester.edu/aincenter/foundry-mentors/>